

Backpropagating Gravitational wave events

KAZE W. K. WONG,¹ KATELYN BREIVIK,¹ AND WILL M. FARR^{1,2}

¹*Center for Computational Astrophysics, Flatiron Institute, New York, NY 10010, USA*

²*Department of Physics and Astronomy, Stony Brook University, Stony Brook NY 11794, USA*

(Dated: May 23, 2022)

ABSTRACT

One promising way to extract information about stellar astrophysics from gravitational wave catalogs is to compare the catalog to the outputs of population synthesis modeling with varying physical assumptions. The parameter space of physical assumptions in population synthesis is high-dimensional and the choice of parameters that best represents the evolution of a binary system may depend in an as-yet-to-be-determined way on the system’s properties; indeed, measuring the variation of population synthesis parameter settings with binary properties is one of the principal ways to extract information about binary astrophysics from catalogs. Here we propose an inference pipeline to simultaneously infer stellar progenitor parameters (e.g. zero-age main sequence masses) and population synthesis parameter settings controlling modeled binary evolution from individual gravitational wave observations of merging compact binaries. Our pipeline can efficiently explore the high-dimensional space of population synthesis settings and progenitor system properties for each system in a catalog of gravitational wave observations. We apply our pipeline to observations in the third LIGO Virgo Kagra Gravitational Wave Transient Catalog (GWTC3). We showcase the effectiveness of this pipeline with a detailed study of the progenitor properties and population synthesis settings that produce mergers like the observed GW150914.



1. INTRODUCTION

As the detection rate of gravitation waves (GWs) from merging double-compact-object (DCO) binaries increases with the sensitivity of the ground-based GW detector network (Abbott et al. 2018; LIGO Scientific Collaboration et al. 2015; Buikema et al. 2020; Tse et al. 2019; Acernese et al. 2015, 2019; Aso et al. 2013; Akutsu et al. 2021), we are beginning to constrain the astrophysical processes which shape the evolution of GW progenitor populations. One of the most common ways to study progenitor populations of GW mergers is through population synthesis simulations for a host of formation environments including isolated binary stars, triple systems, globular clusters, nuclear star clusters, AGN disks,

and more (Mandel & Broekgaarden 2022). WF: Moar citations here Traditionally in these studies, the population is studied through a Monte-Carlo approach. Using theoretically motivated distributions of initial parameters (age, metallicity, masses, orbital separations, eccentricities, etc), initial populations are evolved with fixed astrophysical assumptions (population synthesis parameter settings) to produce a synthetic catalog of mergers which can be compared to observations from GW detectors.

While forward modelling entire populations has been the canonical pathway to compare simulations and GW observations, it is not necessarily the most efficient way to constrain the underlying physical processes, or “hyperparameters”, which drive the evolution of the population. One of the most common skepticisms of population synthesis studies is that they employ simulation sets which often cover a small subdomain of the entire parameter space of uncertain physics, or only vary one uncertain process per simulation (e.g. common envelope efficiency) while holding all others fixed to some

kwong@flatironinstitute.org

kbreivik@flatironinstitute.org

will.farr@stonybrook.edu

“default” values. On top of this, these values are often not calibrated to data. Even with recent developments in using emulators to speed up the simulation process, training the emulators still requires a significant amount of computation to cover a wide range of uncertainty. As an example, to emulate a model with 10 parameters, the simplest way to construct a training set of simulations is to run a grid which varies combinations of uncertain physics. For 10 uncertain physical processes, if we consider 2 variations for each physical process this would still require 1024 simulations. Because population synthesis simulations require hundreds to thousands of cpu hours, training an emulator which spans a high dimensional uncertainty space remains an impossibility at present.

[KW: Reorganizing discussion related to Jeff’s work such that it doesn’t seem like our method is just an extension of DartBoard.]

This compounds with the fact that there is no guarantee the entire population should share the same fixed hyperparameters, as is often assumed in most population synthesis studies. This suggests that there could be a large (possibly infinite) family of hyper-parameter distributions required to compare data to simulations which span the full set of uncertain physical models. However, this issue is not a fundamental issue, but an apparent obstacle which arises from an inefficient approach to the problem. Furthermore, each binary system could in principle come from a different formation environment and hence have different progenitor parameters as well as hyperparameters, which define uncertain physics associated with each environment. In this case, it is advantageous to determine the posterior density in the progenitor and hyper-parameter joint space for each individual binary system, then perform population in analyses in that space. Andrews et al. (2021) (hereafter A21) used DartBoard (Andrews et al. 2018) to determine the ZAMS parameters which produce GW150914-like BBH merger based on posterior samples for GW150914 and a fixed set of hyperparameters which define assumptions for how isolated binary-star interactions proceed using COSMIC, a binary population synthesis code (Breivik et al. 2020). However, the method was limited to a single event instead of the population. More importantly, the method did not vary the hyperparameters while sampling the posterior density of the ZAMS parameters, which limit its capability to be applied to understand a population of binary systems.

In this paper, we propose a method to “backward” model each GW event to its progenitor state while allowing the hyper-parameters to vary across their full range of physical uncertainty following a two-stage process.

First, we solve a root finding problem to obtain guesses that are likely to produce the desired system properties in the joint progenitor–hyper-parameter space. We then sample the posterior in the joint space with Markov chain Monte Carlo (MCMC) algorithm that is initialized using the roots found in the previous stage. The rest of the paper is structured as follows: We detail our method in Sec. 2. In Sec. 3, we show the results of applying our method to a number of real events. We discuss the implications of this work and future direction in Sec. 4.

2. METHOD

We define the progenitor’s zero age main sequence (ZAMS) properties like mass, orbital period, and eccentricity as progenitor parameters θ' , the parameters that control different uncertain physical prescriptions such as wind strength and common envelope efficiency as hyperparameters λ , and random variables that affect certain stochastic process such as whether CO natal kick unbinds a binary as \mathbf{X} . To avoid clutter in the following derivation, we denote all the parameters related to mapping a particular progenitor system into a GW event collectively as evolutionary parameters Θ , which includes θ' , λ , and \mathbf{X} .

Once all the parameters including a particular draw of all the random variables represented by \mathbf{X} are known, the population synthesis code is a deterministic function that transforms the properties of the progenitors into the GW parameters θ :

$$\theta = f(\Theta). \quad (1)$$

The mapping f from Θ to θ can be many-to-one due to degeneracies in the different physical processes and initial parameters of the GW event progenitors. In order to draw inferences about Θ from GW data, we must be able to evaluate the likelihood of that data at fixed progenitor parameters, namely

$$p(d | \Theta) = p(d | \theta = f(\Theta)). \quad (2)$$

The equality holds because the likelihood depends only on the gravitational waveform generated by parameters θ to which the progenitor parameters map onto. In principle this likelihood could be computed at arbitrary values of Θ using the same machinery that is used to estimate source parameters in GW catalogs (Veitch et al. 2015; Ashton et al. 2019; Romero-Shaw et al. 2020; The LIGO Scientific Collaboration et al. 2021) to develop samples of progenitor parameters, hyperparameters, and random variables in Θ from a posterior density.

Since the GW likelihood function ultimately depends only on the GW parameters θ , and GWTC3 already has

samples over these parameters from a posterior density

$$p(\boldsymbol{\theta} | d) \propto p(d | \boldsymbol{\theta}) \pi_{\text{GW}}(\boldsymbol{\theta}), \quad (3)$$

where π_{GW} is the prior density used for sampling, we can save the computation cost associated with evaluating the GW likelihood by rewriting the posterior density of $\boldsymbol{\Theta}$ in terms of the posterior density of $\boldsymbol{\theta}$, namely

$$p(\boldsymbol{\theta} | d) = \frac{p(\boldsymbol{\theta}(\boldsymbol{\Theta}) | d) \pi(\boldsymbol{\Theta})}{\pi_{\text{GW}}(\boldsymbol{\theta}(\boldsymbol{\Theta}))}, \quad (4)$$

which we apply a kernel density estimator to the posterior samples released in GWTC3 to estimate $p(\boldsymbol{\theta}(\boldsymbol{\Theta}) | d)$.

Including hyperparameters and random variables drastically increases the dimensionality of the problem, which makes the sampling process more computationally expensive to converge. To speed up the convergence, we first solve a root-finding problem to find points that are close to the bulk of the posterior density of a GW event, then we use those points to initialize a set of chains in the MCMC process.

For each posterior sample point in the GW-observable space, we can find the corresponding evolutionary parameters by solving an optimization problem that minimizes the mean square difference between the GW event parameters and evolutionary parameters as:

$$\mathcal{L}(\boldsymbol{\Theta}, \boldsymbol{\theta}) = \|\boldsymbol{f}(\boldsymbol{\Theta}) - \boldsymbol{\theta}\|^2. \quad (5)$$

In principle, we should only accept solutions that exactly reproduce the LVK posterior samples in the GW-observable space. However, it is not feasible to achieve such a condition in practice, therefore we relax the condition in eq. 5 to a small acceptance threshold. For this study, we picked a threshold of 10^{-2} . To make sure we find a reasonably complete set of progenitor parameters that corresponds to the posterior sample point, we use 1000 different initial guesses in solving the optimization problem. As long as the solution fulfills the acceptance criteria, the root is as valid as all other roots that fulfill the same acceptance criteria, *regardless to its initial guess*. This means, we have the freedom to choose the set of initial guesses however we think might benefit the optimization.

Most of the points in the evolutionary parameter space do not produce DCO mergers in a Hubble time. For systems that do not merge, we set the binary masses to be 0 such that the gradient for the same systems is also 0

and leads the root finder to get stuck on its first step.¹ Therefore, it is beneficial to choose initial guesses in region of the evolutionary parameter space that are likely to produce DCO mergers. To construct a list of initial guesses, we evolve a set of ZAMS parameters uniformly sampled from the initial binary parameter space, then keep the systems that merge within a Hubble time.

Retaining explicit control over random variables allows us to marginalize over their contribution, and thus allows us to focus on progenitor parameters and hyperparameters. However, in practice most population synthesis codes (including COSMIC which we use in this study) define their random variables implicitly such that the random variables are drawn within the program instead of being passed as arguments. This means the function we use to evolve the binary is not fully deterministic. Thus, even if the root-finding algorithm performs perfectly, forward modelling a set of roots does not guarantee that the simulated population reproduces exactly the set of posterior samples due to randomness in the evolution of each binary. To assure that the recovered progenitor and hyperparameters robustly correspond to the posterior in the GW-observable space, we forward model (or "reproject") the recovered evolutionary parameters to the GW-observable space to check whether the reprojected posterior agrees with the posterior given by the LVK collaboration. We use KL divergence to measure the agreement between the two posterior distributions as:

$$D_{KL}(P||Q) = \int P(\boldsymbol{x}) \log(P(\boldsymbol{x})/Q(\boldsymbol{x})) d\boldsymbol{x}. \quad (6)$$

A small KL divergence means the reprojected COSMIC posterior is similar to the original posterior in the observable space, and is thus a viable channel for that specific event. Otherwise, it either means COSMIC can only explain part of the posterior or cannot explain the event at all. This is completely expected behavior since COSMIC, and isolated binary evolution more generally, carries its own assumptions and is expected to fail in reproducing a subset of the event in GWTC3. One example of this are GW events with at least one component with mass above the pair instability supernova mass limit [cite Farmer, and others](#). The reprojection could also be subject to stochasticity in the evolution of each binary. To account for this, we reproject the posterior multiple times with different random seeds and check whether the KL divergence varies significantly. A

¹ On a plateau of constant loss, the root finder performs no better than a random guess. Because of this, a root finder in a high dimensional space will require a significant amount of computation to get close to a solution.

varying KL divergence means a particular GW event is subject to randomness in the evolution of the progenitor, and extra caution should be used when interpreting the result.

3. RESULTS

We apply the method described in section 2 to the first GW event GW150914. The parameters involved in the analysis and the range we allow for the inference are tabulated in table 1. The spin of BHs formed in isolated binaries depends strongly on the angular momentum transport in massive stars which is still highly uncertain (see however [Bavera+, fuller+](#)), therefore we only consider the two component masses in the observable space in the both the root-finding stage and the MCMC stage. For progenitor parameters, we characterize each guess with five parameters: the two component masses $M_{1,\text{ZAMS}}$ and $M_{2,\text{ZAMS}}$, the orbital period t_{orb} , the eccentricity e , the metallicity Z at ZAMS. Note that the progenitor formation redshift is not obtained through the root-finding process because the formation redshift does not intrinsically affect the evolution of the binary. Instead, we compute the time it takes for the binary to evolve from formation to the merger and add this time back to the lookback time when a GW event is observed and limit the total time to merger to be less than 13.7 Gyr. Using the total lookback time, we can compute the redshift of the binary at ZAMS formation following the Planck 2018 cosmological model ([Planck Collaboration et al. 2020](#)). The main benefit of post-processing the redshift in this way is that we do not rely on a particular assumption of star formation rate (SFR) distribution when we solve for the posterior distribution in redshift. This means we can put the prior on formation redshift (or in general, metallicity-redshift distribution) in by reweighting the posterior distribution according to a particular SFR model, and therefore do not have to rerun the inference everytime we change the SFR model. This is especially helpful since SFR models are still highly uncertain but strongly impact the local merger rates of BBHs ([Broekgaarden et al. 2021](#)). For hyperparameters, we choose parameters that affect binary evolution in COSMIC mostly significantly for massive stars, including the common envelope efficiency α , the fraction of mass accreted onto the accretor during stable mass transfer f_{acc} , the critical mass ratio which determines whether Roche-overflow mass transfer proceeds stably or produces a common envelope when the donor star is on the Hertzsprung gap $q_{\text{crit},3}$, and the root-mean-square of the Maxwellian distribution employed for CO natal kicks, σ .

We show a portion of the joint posterior of the progenitor parameters and hyperparameters of GW150914 in figure 1. Each panel is colored based on the comparison between progenitor parameters (blue), hyperparameters (orange), or a mix of the two (green). We find that the choice for σ does not affect the formation of GW150914-like BBHs. This is not unexpected since we use the ? delayed model which reduces the natal kick strength at BH formation based on the amount of fallback onto the proto-CO. In the case of GW150914-like BBHs, the BH progenitors are massive enough that the fallback reduces the natal kick to zero. Because of this we do not include σ in Figure 1. Similarly, we do not include the ZAMS orbital period and eccentricity, which are correlated with one another, but not strongly with the other progenitor parameters or hyperparameters. We find strong correlations between the ZAMS masses, as well as strong correlations between $M_{2,\text{ZAMS}}$ and f_{acc} , and f_{acc} and α . We show scatter plots of each of these combinations in the upper-right inset of the figure.

The correlations between the ZAMS masses are similar to those found by A21 with the majority of the population preferring primary and secondary masses between 60–90 M_{\odot} . However, we also find some ZAMS masses which extend up to 150 M_{\odot} which is the limit imposed by our assumptions. We explore the formation scenarios which lead to successful GW150914-like mergers and those which fail to produce GW150914-like mergers below in Figure 2.

The correlations between $M_{2,\text{ZAMS}}$ and f_{acc} illustrate the variety of ways which GW150914-like mergers can be produced. For binaries with $M_{2,\text{ZAMS}} < 70 M_{\odot}$, we find that accretion efficiencies of $f_{\text{acc}} > 0.5$ are preferred. One reason for this is based on the requirement that the total mass in the binary must remain above the total mass of GW150914 and strongly non-conservative mass transfer (i.e. $f_{\text{acc}} < 0.5$) reduces the total mass of the system. A compounding factor is that as mass leaves the binary due to non-conservative mass transfer, the evolution of the binary’s orbit is less dramatic and leads to wider binaries on average and thus fewer mergers in a Hubble time. For binaries with $M_{2,\text{ZAMS}} \sim 80 M_{\odot}$, f_{acc} is less constrained. This is because of a preference for these secondaries to also have primary masses near 80 M_{\odot} and thus enter a double-core common envelope evolution in which both stars’ envelopes are ejected, leaving behind two stripped helium cores in a tight orbit. In this case, f_{acc} does not affect the binary evolution and is thus unconstrained. This effect can also be seen in the ZAMS mass and $q_{\text{crit},3}$ panels for masses near 80 M_{\odot} because double-core common-envelope evolution is triggered in COSMIC when the radii of the two stars touch

Parameters	Description	Optimization range
Observables		
$m_{1,\text{GW}}$	Primary mass of the GW event	NA
$m_{2,\text{GW}}$	Secondary mass of the GW event	NA
Progenitor parameters		
$m_{1,\text{ZAMS}}$	Primary mass at ZAMS	$[10, 150] M_{\odot}$
$m_{2,\text{ZAMS}}$	Secondary mass at ZAMS	$[10, 150] M_{\odot}$
t_{orb}	Orbital period at ZAMS	$[5, 5000]$ days
e	Eccentricity at ZAMS	$[0, 1]$
Z	Metallicity at ZAMS	$[10^{-4}, 2 \times 10^{-2}]$
z	Redshift at ZAMS	NA
hyperparameters		
α	Common envelope efficiency	$[0.1, 10]$
f_{acc}	Fraction of accretion during stable mass transfer	$[0, 1]$
$q_{\text{crit},3}$	Critical mass ratio on the Hertzsprung Gap	$[0.1, 10]$
σ	Root mean square of Maxwellian natal kick distribution	$[10, 300]$ km/s

Table 1. A list of parameters we used in this study.

due to the rapid expansion of the primary star upon helium ignition. In this case, the choice for Roche-overflow mass transfer to proceed stably (as prescribed by f_{acc}) or unstably (as prescribed by α) using $q_{\text{crit},3}$ is totally irrelevant, and thus unconstrained. Finally, for binaries with $M_{2,\text{ZAMS}} > 80$, we find that f_{acc} is correlated to decrease with increasing mass. This is due to limits on the total mass which must be ejected from the binary to produce BH masses that match GW150914, the reverse situation to $M_{2,\text{ZAMS}} < 70 M_{\odot}$. Binaries with both component masses near $150 M_{\odot}$ must either go through a common envelope evolution (in which case f_{acc} is unconstrained), or very non-conservative mass transfer (where f_{acc} is low, to produce BBHs with the proper mass.

Finally, we find that α and f_{acc} are largely uncorrelated, though there exist independent trends in each hyperparameter. This is not totally unexpected since the physical processes which are described by each hyperparameter, i.e. stable Roche-lobe overflow and common envelope evolution are two independent channels. Generally, we find that GW150914-like BBHs tend to prefer larger accretion efficiencies and have common envelope ejection efficiencies peaking near $\alpha = 1$. One shortcoming of our method implementation is the application of a single hyperparameter for f_{acc} and α for each binary, while both the primary and the secondary star could in principle be defined by their own hyperparameters that prescribe the outcomes for when each star fills its Roche lobe. We reserve full treatment of this for future work but note that this improvement could reveal cor-

relations between hyperparameters that are not present in our current analysis.

Once we obtain a set of binaries which successfully map the ZAMS parameters and hyperparameters to BBH merger masses, we re-evolve the set of ZAMS parameters with the same physical assumptions as A21, but vary the common envelope efficiency to explore how keeping a fixed model which only varies one hyperparameter contrasts to our results. Figure 2 shows the distribution of ZAMS masses for three models (first through third columns), each with a different α but the same hyper-parameters as A21, as well as the results of our sampling which allows our hyper-parameters to vary (fourth column). We find that holding the accretion efficiency, f_{acc} , and common envelope ejection efficiency, α , to fixed values greatly reduces the ZAMS parameter space that produces GW150914-like mergers. In contrast, by allowing the hyper-parameters to fully span the model uncertainty, we find that there are distinct ZAMS parameters which produce GW150914-like mergers.

Because we can explore the full evolutionary parameter space, we can also determine that there are ZAMS parameters which fail to produce GW150914-like mergers *regardless of our hyper-parameter choice*. For primary ZAMS masses between $\sim 95 - 135 M_{\odot}$ and secondary ZAMS masses between $\sim 50 - 95 M_{\odot}$, we find that there is no combination of accretion efficiency and common envelope ejection efficiency which produces merging BBHs that have masses consistent with GW150914's posteriors. In this region, if BBHs with




Figure 1. The posterior for GW150914 in both the progenitor parameter space and hyper-parameter space. M_1 and M_2 are the progenitors' masses. $\log Z$ is log metallicity at ZAMS. z is the redshift at ZAMS. α is the common envelope efficiency. f_{acc} is the fraction of mass accreted during stable mass transfer. $q_{\text{crit},3}$ is the critical mass ratio on the Hertzsprung Gap, Note that the redshift is not fitted during the root finding process or the MCMC process. Once we have find the evolutionary parameters, we add the delay time to the lookback time of the observed posterior sample, then from the total lookback time we can compute the redshift at ZAMS. We highlight three panels in the corner plots to show the fine structure of the set of posterior samples in the evolutionary parameter space. We also color the posterior in a particular panel according to the type of parameters involved in the corner. Blue denotes panels that include only progenitor parameters, green denotes panels that include a mix of progenitor and hyperparameters, and orange denotes panels that include only hyperparameters.

masses consistent with GW150914’s masses are produced through a combination of α and f_{acc} , they do not merge in a Hubble time. Conversely, in the region of the fourth column of Figure 2 which does produce GW150914-like mergers, the combination of α , f_{acc} , and the ZAMS masses and orbital separations balance to produce BHs with the correct masses that merge within a Hubble time.

WF: What if we made this a subsection of “methods”, and then just one sentence here about the values of the KL divergence for GW150914? To check the validity of the recovered posterior over evolutionary parameters, it is essential to check whether it is consistent with the original posterior in the observable space. We reproject the recovered posterior to the observable space, and compare the KL divergence between the original posterior and the reprojected posterior. In particular, we compute KL divergence of the samples obtained from root-finding alone and after MCMC sampling. WF: Clean up wording. For GW150914, the KL divergence between the reprojected posterior and the original posterior is [KW: fill later] nats on average if we use the root-finding result alone, whereas the MCMC result has a higher KL divergence at [KW: fill] nats. The KL divergence stays stable for reprojecting with different random seeds. In fig.3, we can see the reprojected posterior from both the root-finding process alone and the MCMC result agrees well with the original posterior, even though the agreement between the MCMC result and the original posterior is not perfect. This degradation of the agreement comes from the fact that we are using a KDE to estimate the posterior density in the observable space. Note that the result from root-finding alone and the result from MCMC agrees well in the evolutionary parameter space. Compared to the result shown in a previous study, our result covers almost all the posterior of GW150914 in the masses space, which means our result is more consistent with the observation.

To illustrate the potential benefit of our method on the population level, we perform the same analysis for all events in GWTC3. Plotted in figure 4 is the posterior density in the $m_{1,\text{GW}} - f_{\text{acc}}$ space for most of the events in GWTC3. Each contour represents the 68% credible interval of the posterior density for that particular event. There is a suggestive trend showing f_{acc} could increase as the mass of the progenitor increases. Such a trend implies the stable mass transfer (SMT) phase of a less massive binary would be preferentially non-conservative, with more conservative mass transfer in more massive binary systems.

Note that some events did not pass the KL divergence test we proposed in section 2; we do not include these systems in figure 4. Binaries that form low-mass compact objects have lower amounts of fallback and tend to have correspondingly larger variance in their random natal kicks. Larger kicks can unbind the progenitor binary, or lead to wide binaries that do not merge within a Hubble time. In these cases, the extra variance during the reprojection can produce a posterior that may not agree with the original posterior, hence yielding a higher KL divergence.

For high-mass binaries, COSMIC struggles to produce events above the pair instability supernova mass cutoff $m_{1,\text{GW}}$, so the posterior in the evolutionary parameter space only corresponds to part of the posterior in the observable space below this cutoff, and therefore events beyond the PISN mass gap also have higher KL divergence.

Events with more extreme mass ratios are hard to produce with COSMIC, and therefore less likely to be accurately recovered by our method. This could be due to our method assuming a single value for f_{acc} and α as discussed in Section 3.

Merger observations for which COSMIC struggles are likely to be very informative about formation channels and binary evolution physics, and are therefore likely worthy of close study. We anticipate that future work will follow up these events in detail.

Figure 4 shows that our method could in principle reveal the correlation between progenitor parameters and hyperparameters on a population level. A careful treatment of all events in the catalog and discussion related to the detailed physical implications of figure 4 is beyond the scope of this paper. We defer a detailed study of the physics related to the population of GW events to future work.

4. DISCUSSION

We present a new pathway to understand binary evolution with GW events in this paper. Instead of forward modelling an assumed distribution of initial binaries to the observed population, we find the corresponding evolutionary parameters event-by-event. In this first work, we showcase the power of the proposed method with an application to GW150914. We show the joint posterior of the event’s progenitor parameters and hyperparameters. Our work is both a more efficient way to study GW event progenitors, and also allows the possibility of constraining the astrophysics related to binary evolution, especially by capturing the correlation between hyperparameters in different systems.




Figure 2. Comparison of ZAMS masses for binaries which produce GW150914-like mergers for three variations of α with fixed set of parameter assumptions matching those of A22 (first three columns) and for binaries which produce GW150914-like mergers when α , f_{acc} , and $q_{c,3}$ are allowed to vary, (fourth column).

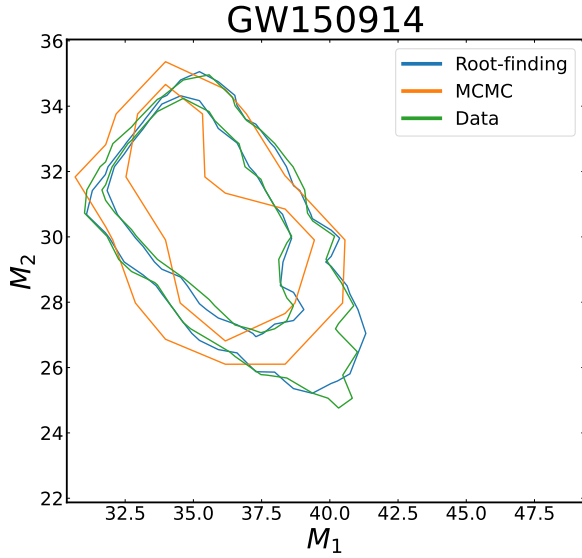


Figure 3. Reprojecting the posterior in evolutionary parameter space of GW150914 to observable space. The blue contour is the reprojected posterior after the root-finding procedure. The orange contour is the reprojected posterior after the MCMC procedure. The green contour is the posterior plotted using the original LVK posterior samples.

Our results suggest that the accretion efficiency during stable mass transfer may depend on the primary black hole mass. In general, our method returns the joint posterior distribution of progenitor parameters and hyperparameters for each event, which enables a data-driven way to study the distribution of hyperparameters. That is, once we have a catalog of events, each has their own posterior in the hyperparameter space, and we can employ well-known techniques such as hierarchical Bayesian analysis to fit a population model to the distribution of hyperparameters; this avoids making overly-specific assumptions like fixing the hyperparameters for all types of event.




Figure 4. The posterior density in the $m_{1,\text{GW}} - f_{\text{acc}}$ space for most of the events in GWTC3. Each contour is the 68% credible interval of the posterior density for a particular event. [KW: Check interval number] At $m_{1,\text{GW}} \sim 45M_{\odot}$, the pair instability supernova mechanism prevents COSMIC from producing events that are more massive than this cutoff. Therefore, events with the majority of posterior support above this cutoff is not compatible with COSMIC, hence has a large KL divergence and excluded from this figure. On the low mass end, neutron star binaries or neutron star-black hole binaries are subjected to randomness induced by the natal kick, also resulting in a larger and fluctuating KL divergence, therefore they are also excluded from the analysis.

Working event-by-event as we do permits post-processing application of arbitrarily complicated models of star formation and metallicity evolution by reweighting the samples we obtain WF: Cite at least Lieke here. Once we have pulled the GW event posterior back from the observable space to the evolutionary parameter space, we can apply the same hierarchical inference methods currently used to infer the compact binary pop-

ulation to our progenitor population. In particular, we can use our progenitor population to infer the formation rate of compact binary population progenitor systems over cosmic time without needing to make artificial and simplistic assumptions about the delay time distribution or the metallicity-specific star formation rate [WF: again, cite Lieke, and also Vitale, Farr, Ng, Rodriguez?](#). Comparing the formation rate of progenitors to the star formation rate provides an avenue to check our understanding of binary evolution and the relative rate of other formation channels.

While our method allows data-driven exploration of the hyperparameters space for the first time, there are a number of improvements can be implemented in future studies. In this study we use only `COSMIC` as our evolutionary function, which by design cannot explain all the events in GWTC3. For example, event with either of the component mass larger than the PISN gap such as GW190521 cannot be explained by `COSMIC`. Alternative channels such as dynamic formation channel will be needed to explain some subset of the events in GWTC3. As pointed out in the literature [\[KW: Cite mike \]](#), it is unlikely that all the observed events in GWTC3 can be explained by a single formation channel. As GW detectors sensitivity increase, we expect to see more and more events that are unusual in some way. Therefore, having a self-consistence population synthesis code that contains multiple formation channels is essential to accommodate the growing catalog of GW events. In this paper, our main focus is to illustrate the concept of "back-propagating" GW events posterior sample in this work, highlighting the capability of our method and motivating the community to build the next generation population synthesis code that can work with our method. To avoid cluttering of focus, we discuss the physical implication of the result presented in this work under our specific assumption (i.e. using `COSMIC` as our evolutionary function) in a companion paper.

Due to the implicit definition of random variables in `COSMIC`, our evolutionary function is stochastic. This introduces significant inefficiency in our root-finding and sampling algorithm. The main stochasticity in `COSMIC` comes from natal kick, which significantly affect the evolution pathway of low mass events such as BNS events. The effect of natal kick is suppressed for heavier mass events due to fallback. This means events with lighter masses are subject to stochasticity of the function, where the sampling process for heavier events behaves as if the evolutionary function we use are deterministic. Due to computational limitation, we only try 1000 different initial guesses per posterior sample in the root-finding process. This means any posterior sample that has a

probability of merging rarer than 1 in 1000 could be missed. Obviously the problems which comes with the randomness can be alleviated by performing more tries per posterior sample, but this is not scalable in practice. On top of limitation in efficiency, some formation channels require explicit control of random variables by construction. For example, in a dynamical formation scenario such as binaries that form in a globular cluster, each binary has some probability of undergoing a multi-body encounter with another member in the cluster. These encounter probability distributions are either studied with direct N-body simulations or semi-analytical methods. In both cases, each member of the cluster is no longer completely independent of the other members, but coupled through the encounter probability distribution. By studying the encounter probability distribution, we can infer the properties of the environment which the binary lives in. This can only be done if we have explicit control over the random variables that characterize the encounter probability distribution.

Another technical note is we use finite differencing to estimate the gradient of the objective function, which could be a significant source of error near transition points in the evolutionary parameter space. Also, finite differencing is increasing inefficiency as we increase the dimensionality of the problem. To improve the accuracy and efficiency in estimating the gradient of objective function, automatic differentiation is a promising feature that modeler should consider incorporating in their population synthesis code in the future.

To summarize, we propose a novel method to recover the posterior samples in the evolutionary parameter space for each GW event. We point out hyperparameters in the usual population synthesis simulation context are not actually parameters related to the population, but parameters about the evolutionary function. This means the binary evolution functions can be constrained on an individual event basis. We "back propagate" the posterior in the observable space to the evolutionary parameter space, thus allowing us to study hyperparameters and its correlation with progenitor parameters in a data-driven manner. Our method makes less assumptions than the traditional forward modelling approach, which often fix the hyperparameters across the entire population. Since we are not limited to the fixed hyperparameters assumption, we can explore the behavior of the hyperparameters across the population much more efficiently. While our work lays down a data analysis pathway to understand the population of GW events, no physics can be learned without a comprehensive physical model. We hope this letter would motivate the community to build the next generation population synthe-

sis codes that have the following properties: first, they should have explicit control over the random variables so marginalizing over random variables can be done more precisely; second, they should be as automatically differentiable as much as possible so exploring the evolutionary parameter space is efficient. Combining this work

and the next generation population synthesis code, we can explore the full parameter space of binary evolution models with the next-generation GW detectors network in the near future.

5. ACKNOWLEDGEMENT

REFERENCES

- Abbott, B. P., Abbott, R., Abbott, T. D., et al. 2018, *Living Reviews in Relativity*, 21, 3, doi: [10.1007/s41114-018-0012-9](https://doi.org/10.1007/s41114-018-0012-9)
- Acernese, F., Agathos, M., Agatsuma, K., et al. 2015, *Classical and Quantum Gravity*, 32, 024001, doi: [10.1088/0264-9381/32/2/024001](https://doi.org/10.1088/0264-9381/32/2/024001)
- Acernese, F., Agathos, M., Aiello, L., et al. 2019, *PhRvL*, 123, 231108, doi: [10.1103/PhysRevLett.123.231108](https://doi.org/10.1103/PhysRevLett.123.231108)
- Akutsu, T., Ando, M., Arai, K., et al. 2021, *Progress of Theoretical and Experimental Physics*, 2021, 05A101, doi: [10.1093/ptep/ptaa125](https://doi.org/10.1093/ptep/ptaa125)
- Andrews, J. J., Cronin, J., Kalogera, V., Berry, C. P. L., & Zezas, A. 2021, *ApJL*, 914, L32, doi: [10.3847/2041-8213/ac00a6](https://doi.org/10.3847/2041-8213/ac00a6)
- Andrews, J. J., Zezas, A., & Fragos, T. 2018, *ApJS*, 237, 1, doi: [10.3847/1538-4365/aaca30](https://doi.org/10.3847/1538-4365/aaca30)
- Ashton, G., Hübner, M., Lasky, P. D., et al. 2019, *ApJS*, 241, 27, doi: [10.3847/1538-4365/ab06fc](https://doi.org/10.3847/1538-4365/ab06fc)
- Aso, Y., Michimura, Y., Somiya, K., et al. 2013, *PhRvD*, 88, 043007, doi: [10.1103/PhysRevD.88.043007](https://doi.org/10.1103/PhysRevD.88.043007)
- Breivik, K., Coughlin, S., Zevin, M., et al. 2020, *ApJ*, 898, 71, doi: [10.3847/1538-4357/ab9d85](https://doi.org/10.3847/1538-4357/ab9d85)
- Broekgaarden, F. S., Berger, E., Stevenson, S., et al. 2021, arXiv e-prints, arXiv:2112.05763, <https://arxiv.org/abs/2112.05763>
- Buikema, A., Cahillane, C., Mansell, G. L., et al. 2020, *PhRvD*, 102, 062003, doi: [10.1103/PhysRevD.102.062003](https://doi.org/10.1103/PhysRevD.102.062003)
- LIGO Scientific Collaboration, Aasi, J., Abbott, B. P., et al. 2015, *Classical and Quantum Gravity*, 32, 074001, doi: [10.1088/0264-9381/32/7/074001](https://doi.org/10.1088/0264-9381/32/7/074001)
- Mandel, I., & Broekgaarden, F. S. 2022, *Living Reviews in Relativity*, 25, 1, doi: [10.1007/s41114-021-00034-3](https://doi.org/10.1007/s41114-021-00034-3)
- Planck Collaboration, Aghanim, N., Akrami, Y., et al. 2020, *A&A*, 641, A6, doi: [10.1051/0004-6361/201833910](https://doi.org/10.1051/0004-6361/201833910)
- Romero-Shaw, I. M., Talbot, C., Biscoveanu, S., et al. 2020, *MNRAS*, 499, 3295, doi: [10.1093/mnras/staa2850](https://doi.org/10.1093/mnras/staa2850)
- The LIGO Scientific Collaboration, the Virgo Collaboration, the KAGRA Collaboration, et al. 2021, arXiv e-prints, arXiv:2111.03606, <https://arxiv.org/abs/2111.03606>
- Tse, M., Yu, H., Kijbunchoo, N., et al. 2019, *PhRvL*, 123, 231107, doi: [10.1103/PhysRevLett.123.231107](https://doi.org/10.1103/PhysRevLett.123.231107)
- Veitch, J., Raymond, V., Farr, B., et al. 2015, *PhRvD*, 91, 042003, doi: [10.1103/PhysRevD.91.042003](https://doi.org/10.1103/PhysRevD.91.042003)